

MCT
MANJARA CHARITABLE TRUST
RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI

Experiment no. 2

Roll no.: 859

Aim: Apply data cleaning techniques (data imputation)

Theory:

A. What is Data Cleaning?

Data cleaning is the process of **detecting, correcting, or removing inaccurate, incomplete, inconsistent, or irrelevant data** from a dataset to improve its quality and reliability for analysis and decision-making.

Dirty data may include:

- Missing values
- Duplicate records
- Incorrect data types
- Outliers
- Inconsistent formatting

High-quality data leads to **accurate analysis, better models, and valid conclusion**

B. DATA CLEANING CYCLE

The data cleaning cycle consists of the following steps:

1. Data Collection

Gather data from different sources such as databases, surveys, or APIs.

2. Data Inspection

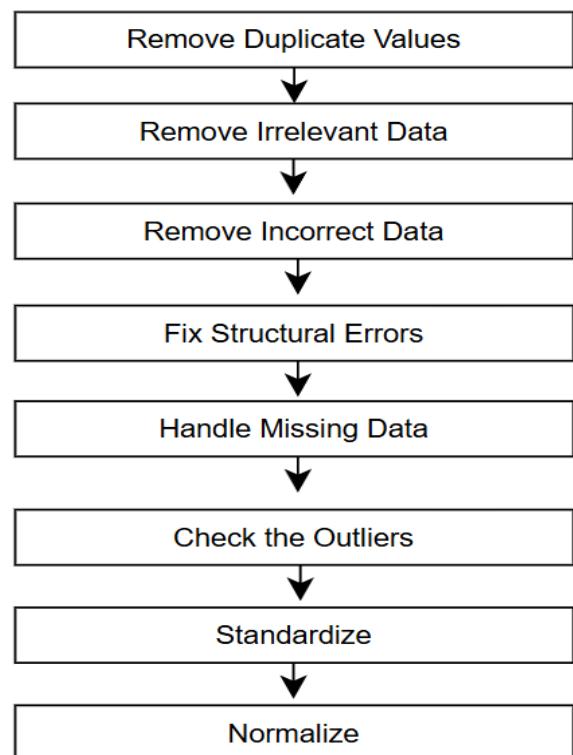
Identify:

- Missing values
- Outliers
- Inconsistencies
- Duplicate entries

3. Data Cleaning

Apply techniques such as:

- Removing duplicates
- Correcting data types
- Handling missing values (imputation)
- Standardizing formats



4. Data Validation

Check whether the cleaned data:

- Meets business rules
- Has logical consistency
- Is complete and accurate

5. Data Storage

Store cleaned data for analysis or modeling.

C. REASONS FOR DATA CLEANING

Data cleaning is necessary because:

- Raw data is often incomplete and inconsistent
- Dirty data leads to wrong insights and poor decisions
- Machine learning models require clean data
- Improves accuracy, reliability, and efficiency
- Saves time during analysis

D. IMPUTATION METHODS

What is Data Imputation?

Data imputation is the process of **replacing missing values with estimated values** rather than removing records.

Common Imputation Methods

- Mean / Median / Mode imputation
- Forward & backward fill
- Regression imputation
- **K-Nearest Neighbors (KNN) imputation**

E. KNN IMPUTATION

What is KNN Imputation?

KNN imputation fills missing values using the **average of the K nearest data points** based on similarity.

How It Works:

1. Select number of neighbors (K)
2. Find K most similar records (distance-based)
3. Compute mean of neighbors
4. Replace missing value

Advantages:

- Uses relationships between variables
- More accurate than mean imputation
- Works well with numerical data

Disadvantages:

- Computationally expensive
- Sensitive to scaling
- Not suitable for very large datasets

Cpde:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

df = sns.load_dataset("titanic")
df.to_csv("orig_titanic_dataset.csv", index=False)
print("Dataset Loaded Successfully")
print(df.head())

print("\nDataset Info:")
print(df.info())

print("\nStatistical Summary:")
print(df.describe())

print("\nMissing Values Count:")
print(df.isnull().sum())

df['age'] = df['age'].fillna(df['age'].median())
df['fare'] = df['fare'].fillna(df['fare'].median())
df['embarked'] = df['embarked'].fillna(df['embarked'].mode()[0])
df['deck'] = df['deck'].fillna(df['deck'].mode()[0])
df['embark_town'] = df['embark_town'].fillna(df['embark_town'].mode()[0])

before = df.shape[0]
df.drop_duplicates(inplace=True)
after = df.shape[0]

print(f"\nDuplicates Removed: {before - after}")

categorical_cols = ['sex', 'embarked', 'class', 'who', 'adult_male', 'alone']

for col in categorical_cols:
    df[col] = df[col].astype(str).str.lower()

df['survived'] = df['survived'].astype(int)
df['pclass'] = df['pclass'].astype(int)

Q1 = df['fare'].quantile(0.25)
Q3 = df['fare'].quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

df['fare'] = np.where(
    df['fare'] > upper_bound,
    upper_bound,
```

```

        np.where(df['fare'] < lower_bound, lower_bound, df['fare'])
    )

print("\nFinal Missing Values:")
print(df.isnull().sum())

print("\nFinal Dataset Info:")
print(df.info())

print("\nCleaned Dataset Preview:")
print(df.head())

df.to_csv("cleaned_titanic_dataset.csv", index=False)
print("\nCleaned dataset saved as 'cleaned_titanic_dataset.csv'")

```

OUTPUT:

| | survived | pclass | sex | age | sibsp | parch | fare | embarked | class | who | adult_male | deck | embark_town | alive | alone |
|---|----------|--------|--------|------|-------|-------|---------|----------|-------|-------|------------|------|-------------|-------|-------|
| 0 | 0 | 3 | male | 22.0 | 1 | 0 | 7.2500 | S | Third | man | True | NaN | Southampton | no | False |
| 1 | 1 | 1 | female | 38.0 | 1 | 0 | 71.2833 | C | First | woman | False | C | Cherbourg | yes | False |
| 2 | 1 | 3 | female | 26.0 | 0 | 0 | 7.9250 | S | Third | woman | False | NaN | Southampton | yes | True |
| 3 | 1 | 1 | female | 35.0 | 1 | 0 | 53.1000 | S | First | woman | False | C | Southampton | yes | False |
| 4 | 0 | 3 | male | 35.0 | 0 | 0 | 8.0500 | S | Third | man | True | NaN | Southampton | no | True |

Dataset Info:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   survived    891 non-null    int64  
 1   pclass      891 non-null    int64  
 2   sex         891 non-null    object 
 3   age         714 non-null    float64 
 4   sibsp       891 non-null    int64  
 5   parch       891 non-null    int64  
 6   fare         891 non-null    float64 
 7   embarked    889 non-null    object 
 8   class        891 non-null    category
 9   who          891 non-null    object 
 10  adult_male  891 non-null    bool   
 11  deck         203 non-null    category
 12  embark_town 889 non-null    object 
 13  alive        891 non-null    object 
 14  alone        891 non-null    bool  
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB
None

```

Final Dataset Info:

```

<class 'pandas.core.frame.DataFrame'>
Index: 779 entries, 0 to 890
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   survived    779 non-null    int64  
 1   pclass      779 non-null    int64  
 2   sex         779 non-null    object 
 3   age         779 non-null    float64 
 4   sibsp       779 non-null    int64  
 5   parch       779 non-null    int64  
 6   fare         779 non-null    float64 
 7   embarked    779 non-null    object 
 8   class        779 non-null    object 
 9   who          779 non-null    object 
 10  adult_male  779 non-null    object 
 11  deck         779 non-null    category
 12  embark_town 779 non-null    object 
 13  alive        779 non-null    object 
 14  alone        779 non-null    object 
dtypes: category(1), float64(2), int64(4), object(8)
memory usage: 92.4+ KB
None

```

Statistical Summary:

| | survived | pclass | age | sibsp | parch | fare |
|-------|------------|------------|------------|------------|------------|------------|
| count | 891.000000 | 891.000000 | 714.000000 | 891.000000 | 891.000000 | 891.000000 |
| mean | 0.383838 | 2.308642 | 29.699118 | 0.523008 | 0.381594 | 32.204208 |
| std | 0.486592 | 0.836071 | 14.526497 | 1.102743 | 0.806057 | 49.693429 |
| min | 0.000000 | 1.000000 | 0.420000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.000000 | 2.000000 | 20.125000 | 0.000000 | 0.000000 | 7.910400 |
| 50% | 0.000000 | 3.000000 | 28.000000 | 0.000000 | 0.000000 | 14.454200 |
| 75% | 1.000000 | 3.000000 | 38.000000 | 1.000000 | 0.000000 | 31.000000 |
| max | 1.000000 | 3.000000 | 80.000000 | 8.000000 | 6.000000 | 512.329200 |

| Missing Values Count: | | Final Missing Values: | |
|-------------------------|-----|-----------------------|---|
| survived | 0 | survived | 0 |
| pclass | 0 | pclass | 0 |
| sex | 0 | sex | 0 |
| age | 177 | age | 0 |
| sibsp | 0 | sibsp | 0 |
| parch | 0 | parch | 0 |
| fare | 0 | fare | 0 |
| embarked | 2 | embarked | 0 |
| class | 0 | class | 0 |
| who | 0 | who | 0 |
| adult_male | 0 | adult_male | 0 |
| deck | 688 | deck | 0 |
| embark_town | 2 | embark_town | 0 |
| alive | 0 | alive | 0 |
| alone | 0 | alone | 0 |
| dtype: int64 | | dtype: int64 | |
| Duplicates Removed: 112 | | | |

| Cleaned Dataset Preview: | | | | | | | | | | | | | | | | |
|--------------------------|----------|--------|--------|------|-------|-------|---------|----------|-------|-------|------------|------|-------------|-------|-------|--|
| | survived | pclass | sex | age | sibsp | parch | fare | embarked | class | who | adult_male | deck | embark_town | alive | alone | |
| 0 | 0 | 3 | male | 22.0 | 1 | 0 | 7.2500 | s | third | man | true | C | Southampton | no | false | |
| 1 | 1 | 1 | female | 38.0 | 1 | 0 | 71.2833 | c | first | woman | false | C | Cherbourg | yes | false | |
| 2 | 1 | 3 | female | 26.0 | 0 | 0 | 7.9250 | s | third | woman | false | C | Southampton | yes | true | |
| 3 | 1 | 1 | female | 35.0 | 1 | 0 | 53.1000 | s | first | woman | false | C | Southampton | yes | false | |
| 4 | 0 | 3 | male | 35.0 | 0 | 0 | 8.0500 | s | third | man | true | C | Southampton | no | true | |

CONCLUSION

Data cleaning is a critical step in data preprocessing.

Among imputation methods, **KNN imputation** is effective because it preserves relationships between variables and provides more accurate estimates compared to simple statistical methods.